

ICP- 8

V V S Murthy Kolla
700729142

Video link:

https://drive.google.com/file/d/1kVIOiSPWbZsA1mRnFyQ46nXVrcth9Y-7/view?usp=share_link

Github link: <https://github.com/murthykolla/ICP-8.git>

LeNet Model:

- For hyperparameter tuning, I have included the learning rate for the Adam optimization.
- Also increased the number of epochs to 50.
- I have also used the early stopping method and monitored the validation accuracy improvement with patience 5.
- Plotted the accuracy and loss plot using the history with subplots method for all the models.
- Plotted the confusion matrix heat map for all the models.
- We can see an increase in the test and validation accuracy from the base model.
- Used batch size of 64 to increase the model's accuracy in the .fit method.

AlexNet Model:

- For hyperparameter tuning, I have increased the dropout rate to 0.5. So that the model can learn from unknown neurons during forward and backpropagation.
- Included the learning rate of 0.001 and Adam optimizer to compile the model.
- Increased the kernel size of the first Conv2D layer.
- I have also used the early stopping method and monitored the validation accuracy improvement with patience 5.
- Used the batch size of 32, which better suits the model, to get better accuracy results.
- Plotted the accuracy and loss plot using the history with subplots method for all the models.
- Plotted the confusion matrix heat map for all the models.
- Here, we can see an increase in the validation accuracy and test accuracy score of 63.95%.
- For data visualization, I tested the prediction on a random image and showed the corresponding image.

VGG16 Model:

- Here, I have used the SGD optimizer with parameters of learning rate of 0.01, decay of 1e-6, momentum 0.9, and Nesterov as True.
- Then I called this optimizer into the compile method to get better accuracy results with SGD optimizer.
- Changed the number of epochs to 25 and batch size to 128 for parameter tuning for better results.
- Then, I saved the model to a .h5 file and plotted the accuracy and loss graphs.
- Plotted the heatmap of the confusion matrix.
- For data visualization, I tested the prediction on a random image and showed the corresponding image.
- Finally, I have loaded the saved .h5 file to predict the test image and printed the prediction class and expected class.

VGG19

- Increased the number of epochs to 100 and batch size to 64.
- Changed the model architecture with the Cifar100 dataset.
- Increased the dense layers to 4096 and the last dense layer to 100.
- Included the dropout layer with a drop rate of 0.5.
- Used the model.fit method with a batch size of 128 and a number of epochs as 25.
- Plotted the accuracy and loss plots.
- Printed the testing accuracy of the model and plotted the confusion accuracy.

```
LeNet&AlexNet.ipynb ☆
File Edit View Insert Runtime Tools Help Changes.will not be saved
+ Code + Text ⚡ Copy to Drive
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import RMSprop, Adam
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.metrics import classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")

(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 4s 0us/step

classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]

y_train = y_train.reshape(-1,)

# Reshape converting 2D to 1D
y_test = y_test.reshape(-1,)
y_train = y_train.reshape(-1,)

# This code normalization
x_train = x_train / 255.0
x_test = x_test / 255.0

x_train.shape

(50000, 32, 32, 3)

from tensorflow.keras import layers, models
lenet = keras.models.Sequential([
    keras.layers.Conv2D(6, kernel_size=5, strides=1, activation='relu', input_shape=(32,32,3), padding='same'), #C1
    keras.layers.AveragePooling2D(2,2), #S1
    keras.layers.Conv2D(16, kernel_size=5, strides=1, activation='relu', padding='same'), #C2
    keras.layers.AveragePooling2D(2,2), #S2
    keras.layers.Conv2D(120, kernel_size=5, strides=1, activation='relu', padding='same'), #C3
    keras.layers.Flatten(), #flatten
    keras.layers.Dense(84, activation='relu'), #F1
    keras.layers.Dense(10, activation='softmax') #Output layer
])
```

```
[ ] lenet.summary()

Model: "sequential_3"
-----  

Layer (type)          Output Shape       Param #
-----  

conv2d_9 (Conv2D)     (None, 32, 32, 6)    456  

average_pooling2d_6 (AveragePooling2D) (None, 16, 16, 6) 0  

conv2d_10 (Conv2D)    (None, 16, 16, 16)   2416  

average_pooling2d_7 (AveragePooling2D) (None, 8, 8, 16) 0  

conv2d_11 (Conv2D)    (None, 8, 8, 120)    48120  

flatten_3 (Flatten)   (None, 7680)         0  

dense_6 (Dense)      (None, 84)           645204  

dense_7 (Dense)      (None, 10)           850  

-----  

Total params: 697,046  

Trainable params: 697,046  

Non-trainable params: 0
```

```
❶ from tensorflow.keras.callbacks import EarlyStopping, LearningRateScheduler

# define early stopping callback
earlystop = EarlyStopping(monitor='val_loss', patience=5)

# compile model
lenet.compile(optimizer=Adam(learning_rate=0.001), loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])

# fit model with callbacks
history = lenet.fit(x_train, y_train, batch_size=64, epochs=50, validation_data=(x_test, y_test), callbacks=[earlystop], verbose=1)
```

❷ Epoch 1/50
782/782 [=====] - 6s 5ms/step - loss: 0.0981 - accuracy: 0.9667 - val_loss: 2.9333 - val_accuracy: 0.6355
Epoch 2/50
782/782 [=====] - 4s 4ms/step - loss: 0.0827 - accuracy: 0.9721 - val_loss: 3.1152 - val_accuracy: 0.6287
Epoch 3/50
782/782 [=====] - 4s 5ms/step - loss: 0.0838 - accuracy: 0.9717 - val_loss: 2.8833 - val_accuracy: 0.6398
Epoch 4/50
782/782 [=====] - 4s 5ms/step - loss: 0.0663 - accuracy: 0.9781 - val_loss: 3.0246 - val_accuracy: 0.6283
Epoch 5/50
782/782 [=====] - 4s 5ms/step - loss: 0.0836 - accuracy: 0.9726 - val_loss: 3.1052 - val_accuracy: 0.6349
Epoch 6/50
782/782 [=====] - 3s 4ms/step - loss: 0.0782 - accuracy: 0.9746 - val_loss: 3.1877 - val_accuracy: 0.6405
Epoch 7/50
782/782 [=====] - 4s 5ms/step - loss: 0.0793 - accuracy: 0.9740 - val_loss: 3.1048 - val_accuracy: 0.6314
Epoch 8/50
782/782 [=====] - 4s 5ms/step - loss: 0.0641 - accuracy: 0.9790 - val_loss: 3.0789 - val_accuracy: 0.6395

```

❶ import matplotlib.pyplot as plt

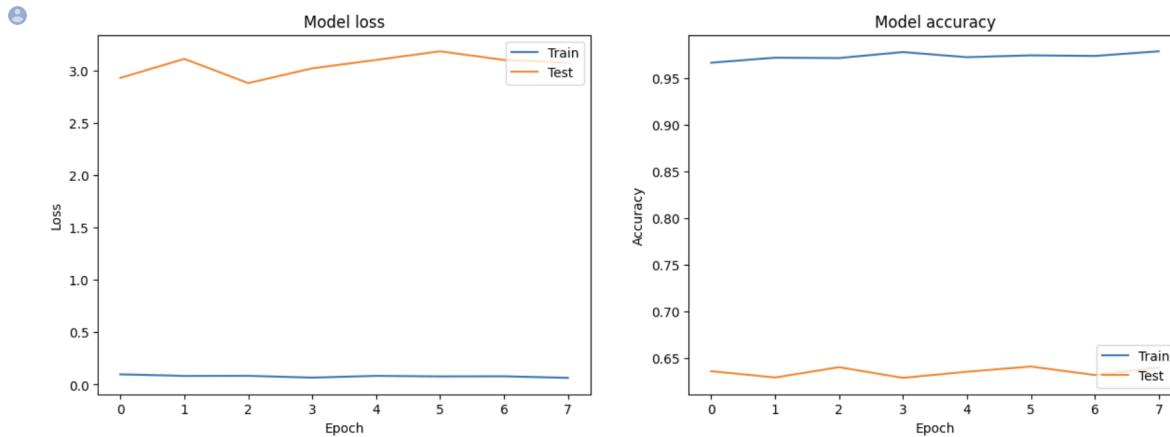
# Plot the results
fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# Plot training and validation accuracy
axs[1].plot(history.history['accuracy'])
axs[1].plot(history.history['val_accuracy'])
axs[1].set_title('Model accuracy')
axs[1].set_ylabel('Accuracy')
axs[1].set_xlabel('Epoch')
axs[1].legend(['Train', 'Test'], loc='lower right')

# Plot training and validation loss
axs[0].plot(history.history['loss'])
axs[0].plot(history.history['val_loss'])
axs[0].set_title('Model loss')
axs[0].set_ylabel('Loss')
axs[0].set_xlabel('Epoch')
axs[0].legend(['Train', 'Test'], loc='upper right')

plt.show()

```



```

[ ] from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1, )
y_predictions= np.argmax(y_predictions, axis=1)

confusion_matrix(y_test, y_predictions)

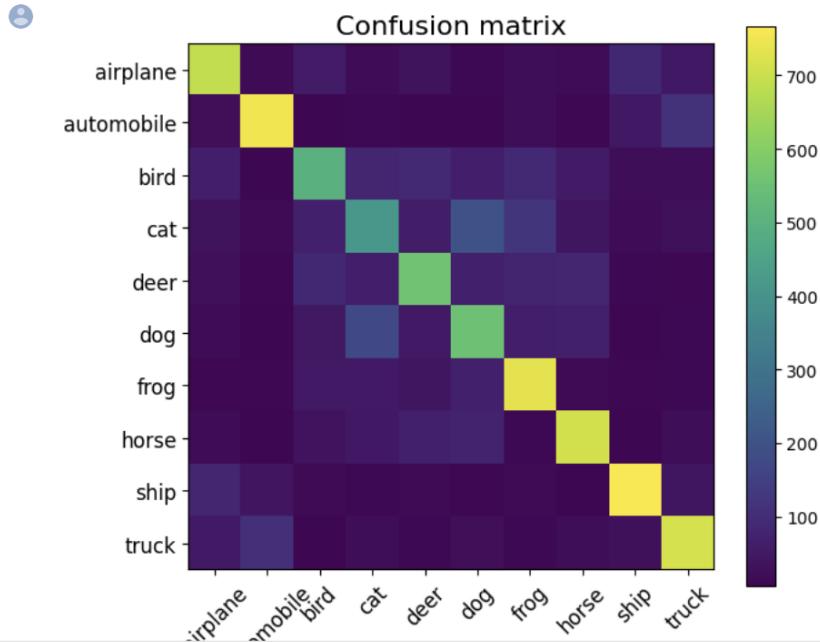
```

```
[ ] from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1,)
y_predictions= np.argmax(y_predictions, axis=1)

confusion_matrix(y_test, y_predictions)
```

313/313 [=====] - 1s 2ms/step
array([[690, 15, 57, 19, 34, 13, 20, 17, 86, 49],
[25, 753, 6, 13, 7, 7, 21, 8, 47, 113],
[67, 7, 501, 82, 92, 67, 89, 54, 20, 21],
[34, 16, 69, 411, 60, 199, 118, 43, 19, 31],
[26, 8, 88, 66, 560, 69, 79, 82, 12, 10],
[17, 7, 46, 173, 52, 553, 65, 70, 5, 12],
[9, 8, 52, 50, 41, 68, 736, 14, 10, 12],
[18, 6, 37, 47, 68, 76, 13, 710, 5, 20],
[81, 38, 14, 13, 16, 10, 14, 5, 766, 43],
[51, 109, 7, 21, 11, 24, 11, 20, 31, 715]])

```
▶ # confusion matrix and accuracy
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



```
[ ] print("Test accuracy:", accuracy_score(y_test, y_predictions))  
Test accuracy: 0.6395
```

```
❸ L = 8  
W = 8  
fig, axes = plt.subplots(L, W, figsize = (20,20))  
axes = axes.ravel() #  
  
for i in np.arange(0, L * W):  
    axes[i].imshow(x_test[i])  
    axes[i].set_title("Predicted = {}\nActual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))  
    axes[i].axis('off')  
  
plt.subplots_adjust(wspace=1)
```

Predicted = cat
Actual = cat



Predicted = ship
Actual = ship



Predicted = ship
Actual = ship



Predicted = airplane
Actual = airplane



Predicted = frog
Actual = frog



Predicted = frog
Actual = frog



Predicted = frog
Actual = automobile



Predicted = cat
Actual = frog



Predicted = cat
Actual = cat



Predicted = automobile
Actual = automobile



Predicted = deer
Actual = airplane



Predicted = truck
Actual = truck



Predicted = dog
Actual = dog



Predicted = horse
Actual = horse



Predicted = automobile
Actual = truck



Predicted = frog
Actual = ship



Predicted = dog
Actual = dog



Predicted = deer
Actual = horse



Predicted = ship
Actual = ship



Predicted = frog
Actual = frog



Predicted = deer
Actual = horse



Predicted = airplane
Actual = airplane



Predicted = deer
Actual = deer



Predicted = truck
Actual = truck



Predicted = deer
Actual = dog



Predicted = deer
Actual = bird



Predicted = horse
Actual = deer



Predicted = airplane
Actual = airplane



Predicted = truck
Actual = truck



Predicted = frog
Actual = frog



Predicted = frog
Actual = frog



Predicted = deer
Actual = dog



```

dropout_24 (Dropout)      (None, 4096)          0
dense_45 (Dense)         (None, 4096)          16781312
dropout_25 (Dropout)      (None, 4096)          0
dense_46 (Dense)         (None, 10)            40970
=====
Total params: 16,992,292
Trainable params: 16,992,292
Non-trainable params: 0

```

```

from tensorflow.keras.callbacks import EarlyStopping, LearningRateScheduler

# define early stopping callback
earlystop = EarlyStopping(monitor='val_loss', patience=5)

# compile model
AlexNet.compile(optimizer=Adam(learning_rate=0.001), loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])

# fit model with callbacks
history = AlexNet.fit(x_train, y_train, batch_size=32, epochs=50, validation_data=(x_test, y_test), callbacks=[earlystop], verbose=1)

Epoch 1/50
1563/1563 [=====] - 13s 6ms/step - loss: 2.0881 - accuracy: 0.1795 - val_loss: 1.8452 - val_accuracy: 0.2964
Epoch 2/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.7712 - accuracy: 0.3294 - val_loss: 1.6706 - val_accuracy: 0.3578
Epoch 3/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.6495 - accuracy: 0.3915 - val_loss: 1.5894 - val_accuracy: 0.4165
Epoch 4/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.5902 - accuracy: 0.4226 - val_loss: 1.5480 - val_accuracy: 0.4264
Epoch 5/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.5533 - accuracy: 0.4366 - val_loss: 1.5612 - val_accuracy: 0.4386
Epoch 6/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.5283 - accuracy: 0.4467 - val_loss: 1.5326 - val_accuracy: 0.4420
Epoch 7/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.5043 - accuracy: 0.4584 - val_loss: 1.5107 - val_accuracy: 0.4561
Epoch 8/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.4799 - accuracy: 0.4726 - val_loss: 1.4988 - val_accuracy: 0.4596
Epoch 9/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.4606 - accuracy: 0.4769 - val_loss: 1.4803 - val_accuracy: 0.4657
Epoch 10/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.4381 - accuracy: 0.4872 - val_loss: 1.4988 - val_accuracy: 0.4687
Epoch 11/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.4211 - accuracy: 0.4955 - val_loss: 1.4805 - val_accuracy: 0.4731
Epoch 12/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.4070 - accuracy: 0.5003 - val_loss: 1.4811 - val_accuracy: 0.4651
Epoch 13/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.3889 - accuracy: 0.5078 - val_loss: 1.4747 - val_accuracy: 0.4763
Epoch 14/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.3821 - accuracy: 0.5123 - val_loss: 1.4708 - val_accuracy: 0.4830
Epoch 15/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.3658 - accuracy: 0.5201 - val_loss: 1.5083 - val_accuracy: 0.4631
Epoch 16/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.3583 - accuracy: 0.5221 - val_loss: 1.4588 - val_accuracy: 0.4778
Epoch 17/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.3425 - accuracy: 0.5276 - val_loss: 1.4721 - val_accuracy: 0.4818
Epoch 18/50
1563/1563 [=====] - 9s 6ms/step - loss: 1.3301 - accuracy: 0.5328 - val_loss: 1.4761 - val_accuracy: 0.4769

```

```
[ ] from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Flatten
from tensorflow.keras.optimizers import SGD, Adam
from keras.layers.convolutional import Convolution2D
from keras.layers.convolutional import MaxPooling2D

➊ #Define Alexnet Model
AlexNet = Sequential()
AlexNet.add(Conv2D(filters=16,kernel_size=(11,11),strides=(4,4),input_shape=(32,32,3), activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(5,5),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(1,1)))
AlexNet.add(Conv2D(60,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(30,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(20,(3,3),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Flatten())
AlexNet.add(Dense(4096, activation='relu'))
AlexNet.add(Dropout(0.5))
AlexNet.add(Dense(4096, activation='relu'))
AlexNet.add(Dropout(0.5))
AlexNet.add(Dense(10,activation='softmax'))

AlexNet.compile(optimizer=Adam(learning_rate=0.001), loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
AlexNet.summary()
```

➌ Model: "sequential_22"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_102 (Conv2D)	(None, 6, 6, 16)	5824
max_pooling2d_54 (MaxPoolin g2D)	(None, 3, 3, 16)	0
conv2d_103 (Conv2D)	(None, 3, 3, 60)	24060
max_pooling2d_55 (MaxPoolin g2D)	(None, 2, 2, 60)	0
conv2d_104 (Conv2D)	(None, 2, 2, 60)	32460
conv2d_105 (Conv2D)	(None, 2, 2, 30)	16230
conv2d_106 (Conv2D)	(None, 2, 2, 20)	5420
max_pooling2d_56 (MaxPoolin g2D)	(None, 1, 1, 20)	0
flatten_16 (Flatten)	(None, 20)	0
dense_44 (Dense)	(None, 4096)	86016
dropout_24 (Dropout)	(None, 4096)	0
dense_45 (Dense)	(None, 4096)	16781312

```

l ] from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Flatten
from tensorflow.keras.optimizers import SGD, Adam
from keras.layers.convolutional import Convolution2D
from keras.layers.convolutional import MaxPooling2D

```

❶ #Define Alexnet Model
AlexNet = Sequential()
AlexNet.add(Conv2D(filters=16,kernel_size=(11,11),strides=(4,4),input_shape=(32,32,3), activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Conv2D(60,(5,5),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(1,1)))
AlexNet.add(Conv2D(60,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(30,(3,3),padding='same',activation='relu'))
AlexNet.add(Conv2D(20,(3,3),padding='same',activation='relu'))
AlexNet.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
AlexNet.add(Flatten())
AlexNet.add(Dense(4096, activation='relu'))
AlexNet.add(Dropout(0.5))
AlexNet.add(Dense(4096, activation='relu'))
AlexNet.add(Dropout(0.5))
AlexNet.add(Dense(10,activation='softmax'))

AlexNet.compile(optimizer=Adam(learning_rate=0.001), loss=keras.losses.sparse_categorical_crossentropy, metrics=['accuracy'])
AlexNet.summary()

❷ Model: "sequential_22"

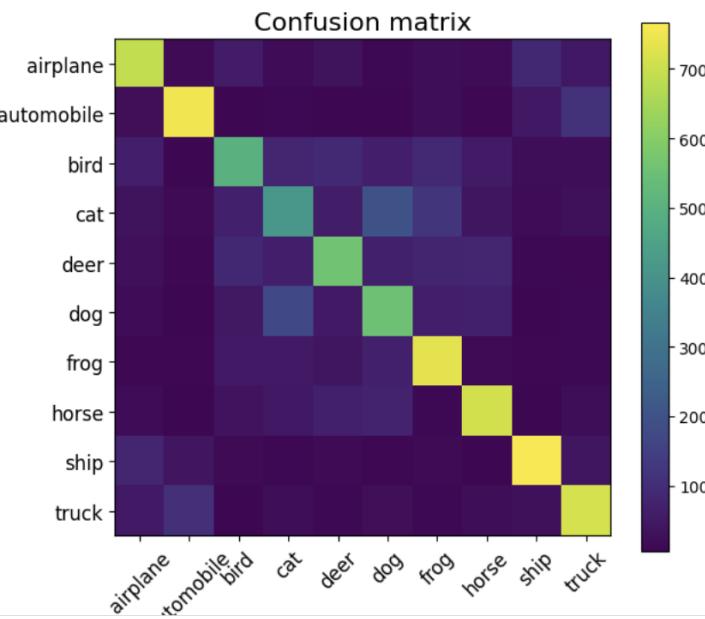
Layer (type)	Output Shape	Param #
<hr/>		
conv2d_102 (Conv2D)	(None, 6, 6, 16)	5824
max_pooling2d_54 (MaxPoolin g2D)	(None, 3, 3, 16)	0
conv2d_103 (Conv2D)	(None, 3, 3, 60)	24060
max_pooling2d_55 (MaxPoolin g2D)	(None, 2, 2, 60)	0
conv2d_104 (Conv2D)	(None, 2, 2, 60)	32460
conv2d_105 (Conv2D)	(None, 2, 2, 30)	16230
conv2d_106 (Conv2D)	(None, 2, 2, 20)	5420
max_pooling2d_56 (MaxPoolin g2D)	(None, 1, 1, 20)	0
flatten_16 (Flatten)	(None, 20)	0
dense_44 (Dense)	(None, 4096)	86016
dropout_24 (Dropout)	(None, 4096)	0
dense_45 (Dense)	(None, 4096)	16781312

```
[ ] from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
y_predictions= lenet.predict(x_test)
y_predictions.reshape(-1, )
y_predictions= np.argmax(y_predictions, axis=1)

confusion_matrix(y_test, y_predictions)

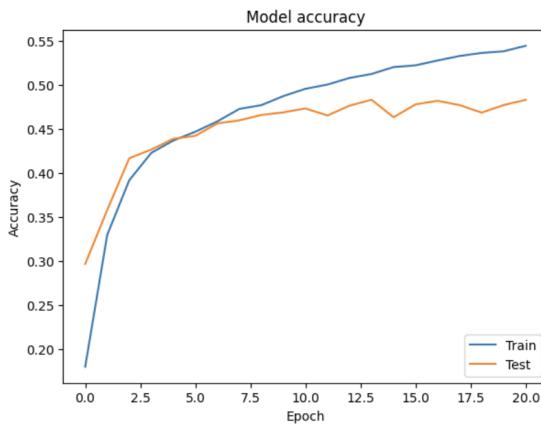
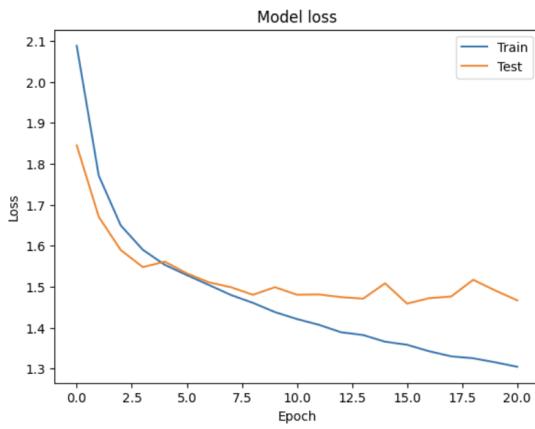
313/313 [=====] - 1s 2ms/step
array([[690, 15, 57, 19, 34, 13, 20, 17, 86, 49],
       [25, 753, 6, 13, 7, 7, 21, 8, 47, 113],
       [67, 7, 501, 82, 92, 67, 89, 54, 20, 21],
       [34, 16, 69, 411, 60, 199, 118, 43, 19, 31],
       [26, 8, 88, 66, 560, 69, 79, 82, 12, 10],
       [17, 7, 46, 173, 52, 553, 65, 70, 5, 12],
       [9, 8, 52, 50, 41, 68, 736, 14, 10, 12],
       [18, 6, 37, 47, 68, 76, 13, 710, 5, 20],
       [81, 38, 14, 13, 16, 10, 14, 5, 766, 43],
       [51, 109, 7, 21, 11, 24, 11, 20, 31, 715]])
```

```
# confusion matrix and accuracy
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



```
[ ] Epoch 20/50  
1563/1563 [=====] - 9s 6ms/step - loss: 1.3155 - accuracy: 0.5381 - val_loss: 1.4908 - val_accuracy: 0.4770  
Epoch 21/50  
1563/1563 [=====] - 9s 6ms/step - loss: 1.3046 - accuracy: 0.5444 - val_loss: 1.4668 - val_accuracy: 0.4829
```

```
▶ import matplotlib.pyplot as plt  
  
# Plot the results  
fig, axs = plt.subplots(1, 2, figsize=(15, 5))  
  
# Plot training and validation accuracy  
axs[1].plot(history.history['accuracy'])  
axs[1].plot(history.history['val_accuracy'])  
axs[1].set_title('Model accuracy')  
axs[1].set_ylabel('Accuracy')  
axs[1].set_xlabel('Epoch')  
axs[1].legend(['Train', 'Test'], loc='lower right')  
  
# Plot training and validation loss  
axs[0].plot(history.history['loss'])  
axs[0].plot(history.history['val_loss'])  
axs[0].set_title('Model loss')  
axs[0].set_ylabel('Loss')  
axs[0].set_xlabel('Epoch')  
axs[0].legend(['Train', 'Test'], loc='upper right')  
  
plt.show()
```



```
[ ] y_predictions1 = AlexNet.predict(x_test)  
y_predictions1.reshape(-1,)
```

```
[ ] print("Test accuracy:", accuracy_score(y_test, y_predictions))  
Test accuracy: 0.6395
```

```
L = 8  
W = 8  
fig, axes = plt.subplots(L, W, figsize = (20,20))  
axes = axes.ravel() #  
  
for i in np.arange(0, L * W):  
    axes[i].imshow(x_test[i])  
    axes[i].set_title("Predicted = {}\nActual = {}".format(classes[y_predictions[i]], classes[y_test[i]]))  
    axes[i].axis('off')  
  
plt.subplots_adjust(wspace=1)
```

Predicted = cat Actual = cat 	Predicted = ship Actual = ship 	Predicted = ship Actual = ship 	Predicted = airplane Actual = airplane 	Predicted = frog Actual = frog 	Predicted = frog Actual = frog 	Predicted = frog Actual = automobile 	Predicted = cat Actual = frog 
Predicted = cat Actual = cat 	Predicted = automobile Actual = automobile 	Predicted = deer Actual = airplane 	Predicted = truck Actual = truck 	Predicted = dog Actual = dog 	Predicted = horse Actual = horse 	Predicted = automobile Actual = truck 	Predicted = frog Actual = ship 
Predicted = dog Actual = dog 	Predicted = deer Actual = horse 	Predicted = ship Actual = ship 	Predicted = frog Actual = frog 	Predicted = deer Actual = horse 	Predicted = airplane Actual = airplane 	Predicted = deer Actual = deer 	Predicted = truck Actual = truck 
Predicted = deer Actual = dog 	Predicted = deer Actual = bird 	Predicted = horse Actual = deer 	Predicted = airplane Actual = airplane 	Predicted = truck Actual = truck 	Predicted = frog Actual = frog 	Predicted = frog Actual = frog 	Predicted = deer Actual = dog 

```

import matplotlib.pyplot as plt

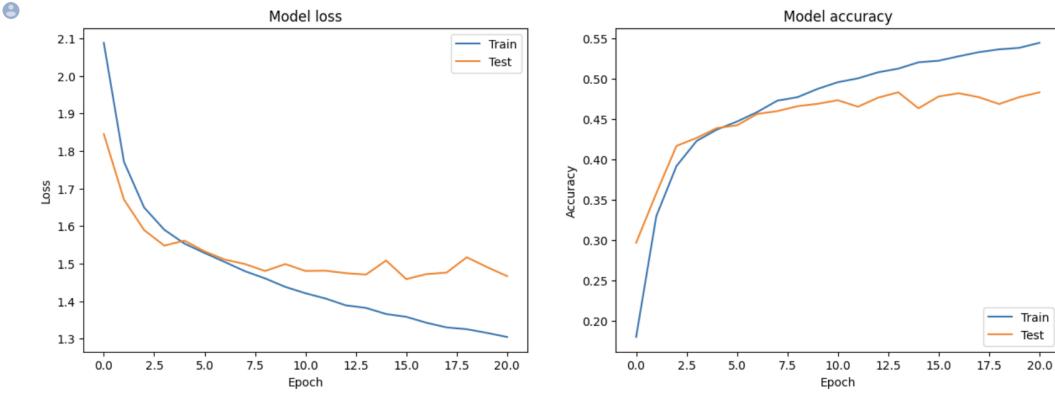
# Plot the results
fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# Plot training and validation accuracy
axs[1].plot(history.history['accuracy'])
axs[1].plot(history.history['val_accuracy'])
axs[1].set_title('Model accuracy')
axs[1].set_ylabel('Accuracy')
axs[1].set_xlabel('Epoch')
axs[1].legend(['Train', 'Test'], loc='lower right')

# Plot training and validation loss
axs[0].plot(history.history['loss'])
axs[0].plot(history.history['val_loss'])
axs[0].set_title('Model loss')
axs[0].set_ylabel('Loss')
axs[0].set_xlabel('Epoch')
axs[0].legend(['Train', 'Test'], loc='upper right')

plt.show()

```



```

y_predictions1 = AlexNet.predict(x_test)
y_predictions1.reshape(-1, )
y_predictions1= np.argmax(y_predictions1, axis=1)

confusion_matrix(y_test, y_predictions1)

313/313 [=====] - 1s 2ms/step
array([[468, 30, 50, 36, 38, 5, 24, 29, 234, 861,

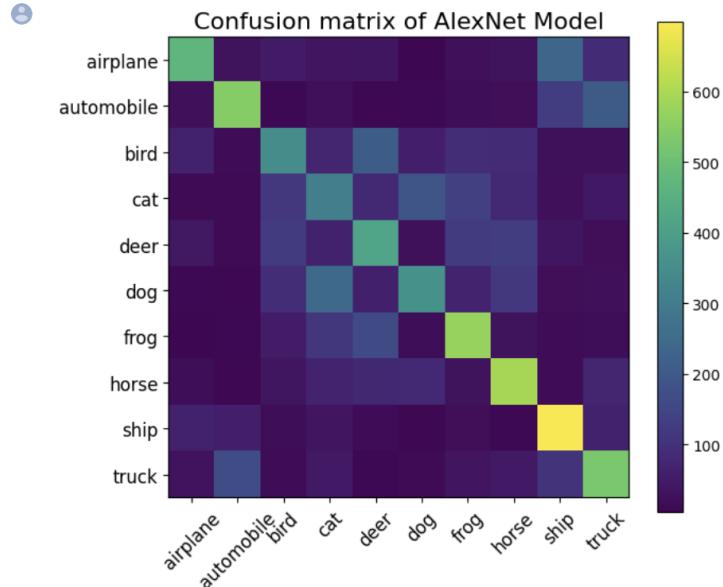
```

```
[ ] y_predictions1 = AlexNet.predict(x_test)
y_predictions1.reshape(-1, )
y_predictions1= np.argmax(y_predictions1, axis=1)

confusion_matrix(y_test, y_predictions1)

313/313 [=====] - 1s 2ms/step
array([[468, 30, 50, 36, 38, 5, 24, 29, 234, 86],
       [23, 545, 10, 24, 9, 10, 20, 22, 130, 207],
       [64, 16, 348, 73, 208, 60, 92, 87, 26, 26],
       [13, 13, 113, 305, 81, 191, 136, 81, 25, 42],
       [41, 13, 125, 65, 415, 28, 125, 130, 37, 21],
       [11, 8, 92, 243, 61, 360, 67, 113, 21, 24],
       [4, 10, 52, 116, 160, 20, 571, 31, 16, 20],
       [18, 9, 38, 68, 75, 80, 29, 594, 15, 74],
       [66, 60, 18, 36, 16, 9, 21, 11, 698, 65],
       [33, 166, 16, 46, 11, 13, 36, 47, 107, 525]])
```

```
▶ # confusion matrix and accuracy
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix of AlexNet Model', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions1))
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



```
▶ # Define the class names for CIFAR-10
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# Choose a random image from the test set
index = np.random.randint(0, x_test.shape[0])
image = x_test[index]

# Make a prediction on the image using the pre-trained model
prediction = AlexNet.predict(np.expand_dims(image, axis=0))

class_index = tf.argmax(prediction, axis=1)
class_name = tf.keras.utils.to_categorical(class_index, num_classes=10)
actual_class = y_test[index]

# Visualize the image and the predicted and actual classes
plt.imshow(image)
plt.title(f"Predicted class")
plt.axis('off')
plt.show()
```

1/1 [=====] - 0s 24ms/step

Predicted class



CO VGG16.ipynb ☆

File Edit View Insert Runtime Tools Help Changes will not be saved

+ Code + Text Copy to Drive

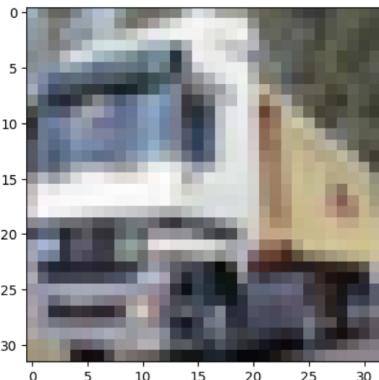
```
[ ] import keras
from keras.models import Sequential
from keras.layers import Activation,Dense,Dropout,Conv2D,Flatten,MaxPooling2D
from keras.datasets import cifar10
from keras import optimizers
from keras.optimizers import SGD
from matplotlib import pyplot as plt

[ ] # generate cifar10 data
(x_train,y_train),(x_test,y_test) = cifar10.load_data()

[ ] # config parameters
num_classes = 10
input_shape = x_train.shape[1:4]
# Compile the model
sgd = SGD(lr=0.01, decay=1e-6, momentum=0.9, nesterov=True)

[ ] # convert label to one-hot
one_hot_y_train = keras.utils.to_categorical(y_train,num_classes=num_classes)
one_hot_y_test = keras.utils.to_categorical(y_test,num_classes=num_classes)

[ ] # check data
plt.imshow(x_train[1])
print(x_train[1].shape)

(32, 32, 3)

[ ] # build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
```

+ Code + Text

```
# build model(similar to VGG16, only change the input and output shape)
model = Sequential()
model.add(Conv2D(64,(3,3),activation='relu',input_shape=input_shape,padding='same'))
model.add(Conv2D(64,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(Conv2D(128,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(Conv2D(256,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(Conv2D(512,(3,3),activation='relu',padding='same'))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Flatten())
model.add(Dense(4096,activation='relu'))
model.add(Dense(4096,activation='relu'))
model.add(Dense(num_classes))
model.add(Activation('softmax'))
```

```
[ ] model.compile(optimizer=sgd, loss='categorical_crossentropy', metrics=['accuracy'])
```

```
# check model
model.summary()
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_52 (Conv2D)	(None, 32, 32, 64)	1792
conv2d_53 (Conv2D)	(None, 32, 32, 64)	36928
max_pooling2d_20 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_54 (Conv2D)	(None, 16, 16, 128)	73856
conv2d_55 (Conv2D)	(None, 16, 16, 128)	147584
max_pooling2d_21 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_56 (Conv2D)	(None, 8, 8, 256)	295168
conv2d_57 (Conv2D)	(None, 8, 8, 256)	590080
conv2d_58 (Conv2D)	(None, 8, 8, 256)	590080

```

import matplotlib.pyplot as plt

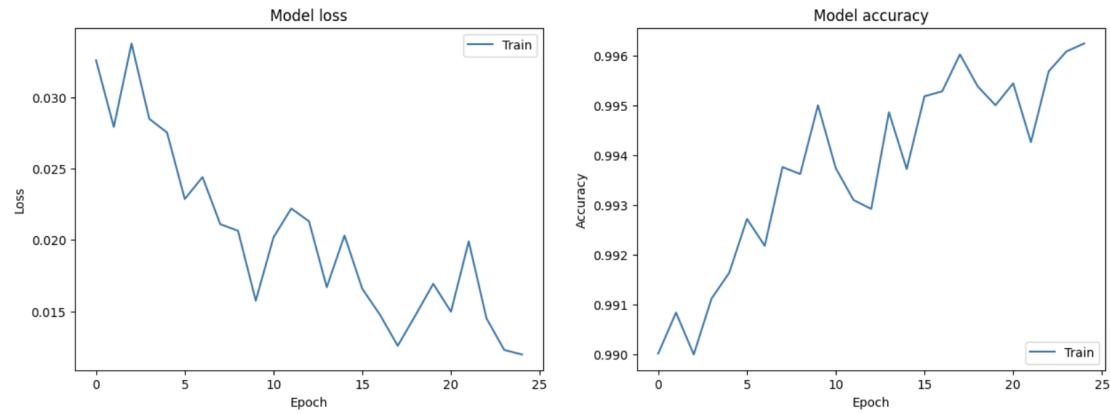
# Plot the results
fig, axs = plt.subplots(1, 2, figsize=(15, 5))

# Plot training and validation accuracy
axs[1].plot(history.history['accuracy'])
axs[1].set_title('Model accuracy')
axs[1].set_ylabel('Accuracy')
axs[1].set_xlabel('Epoch')
axs[1].legend(['Train', 'Test'], loc='lower right')

# Plot training and validation loss
axs[0].plot(history.history['loss'])
axs[0].set_title('Model loss')
axs[0].set_ylabel('Loss')
axs[0].set_xlabel('Epoch')
axs[0].legend(['Train', 'Test'], loc='upper right')

plt.show()

```



```

[ ] import numpy as np
y_predictions1 = model.predict(x_test)
y_predictions1.reshape(-1, )

```

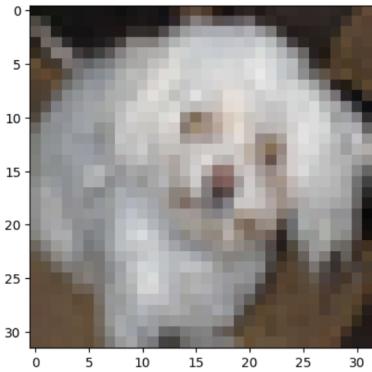
```
[ ] # evaluate
print(model.metrics_names)
model.evaluate(x=x_test,y=one_hot_y_test,batch_size=512)

['loss', 'accuracy']
20/20 [=====] - 0s 8ms/step - loss: 1.5592 - accuracy: 0.7990
[1.5591590404510498, 0.7990000247955322]
```

```
❸ # predict
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i,_ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
print("predict class:",predict)
print("expected class:",expect)
```

```
❹ 1/1 [=====] - 0s 28ms/step
predict class: 5
expected class: 5
```

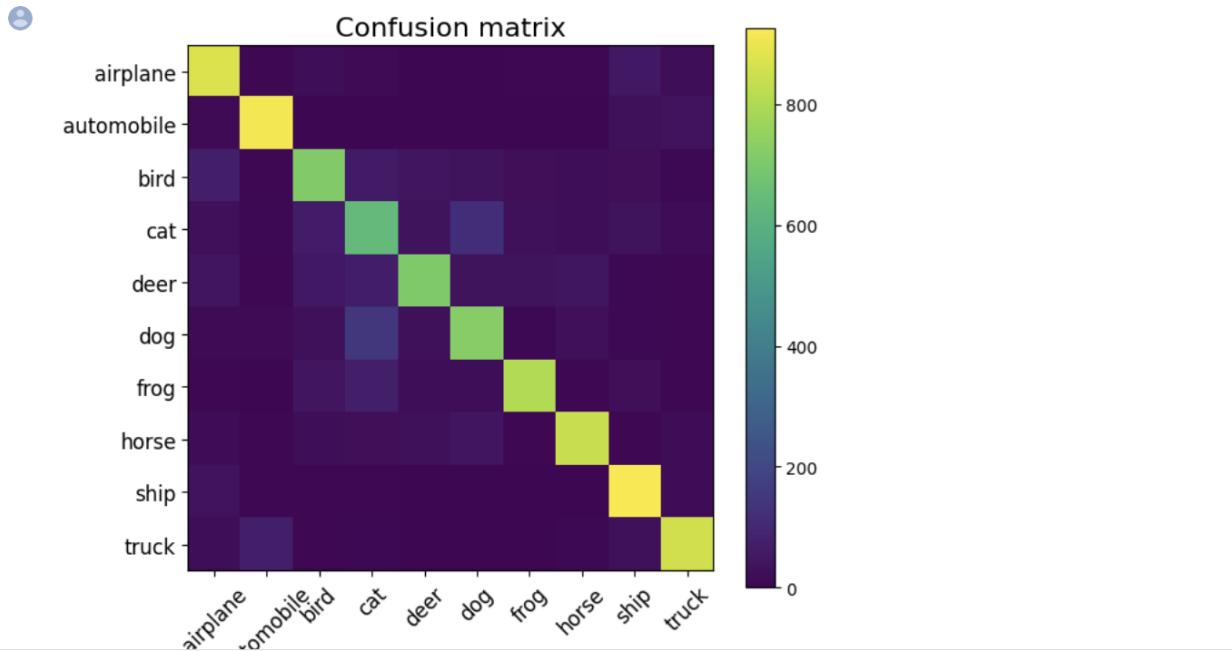


```
[ ] # save model
model.save("keras-VGG16-cifar10.h5")
```

```
[ ] import numpy as np
y_predictions1 = model.predict(x_test)
y_predictions1.reshape(-1, )
y_predictions1= np.argmax(y_predictions1, axis=1)
from sklearn.metrics import classification_report, confusion_matrix
confusion_matrix(y_test, y_predictions1)

313/313 [=====] - 1s 3ms/step
array([[874,    7,   20,   11,    2,    2,    5,    6,   53,   20],
       [ 12,  914,    1,    0,    0,    1,    1,    0,   32,   39],
       [ 71,    5,  714,   58,   41,   35,   22,   19,   25,   10],
       [ 28,   10,   65,  636,   36,  122,   30,   21,   35,   17],
       [ 42,    7,   51,   66,  706,   34,   35,   44,   10,    5],
       [ 12,   11,   32,  141,   31,  723,    8,   28,    8,    6],
       [  7,    3,   40,   72,   19,   20,  803,    6,   25,    5],
       [ 17,    4,   21,   24,   29,   40,    4,  839,    7,   15],
       [ 37,    5,    5,    4,    2,    2,    2,    3,  925,   15],
       [ 21,   70,    4,    8,    0,    1,    3,    8,   29,  856]])
```

```
▶ # confusion matrix and accuracy
from sklearn.metrics import confusion_matrix, accuracy_score
plt.figure(figsize=(7, 6))
plt.title('Confusion matrix', fontsize=16)
plt.imshow(confusion_matrix(y_test, y_predictions1))
classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
plt.xticks(np.arange(10), classes, rotation=45, fontsize=12)
plt.yticks(np.arange(10), classes, fontsize=12)
plt.colorbar()
plt.show()
```



```
[ ] # Define the class names for CIFAR-10
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# Choose a random image from the test set
index = np.random.randint(0, x_test.shape[0])
image = x_test[index]

# Make a prediction on the image using the pre-trained model
prediction = model.predict(np.expand_dims(image, axis=0))
import tensorflow as tf

class_index = tf.argmax(prediction, axis=1)
class_name = tf.keras.utils.to_categorical(class_index, num_classes=10)
actual_class = y_test[index]

# Visualize the image and the predicted and actual classes
plt.imshow(image)
plt.title(f"Predicted class")
plt.axis('off')
plt.show()
```

1/1 [=====] - 0s 28ms/step

Predicted class



```
[ ] from keras.models import load_model
model = load_model('keras-VGG16-cifar10.h5')
# predict
plt.imshow(x_test[1000])

result = model.predict(x_test[1000:1001]).tolist()
predict = 0
expect = y_test[1000][0]
for i,_ in enumerate(result[0]):
    if result[0][i] > result[0][predict]:
        predict = i
```